

# Continuous\_Control

November 24, 2019

## 1 Continuous Control

---

In this notebook, you will learn how to use the Unity ML-Agents environment for the second project of the [Deep Reinforcement Learning Nanodegree](#) program.

### 1.0.1 1. Start the Environment

We begin by importing the necessary packages. If the code cell below returns an error, please revisit the project instructions to double-check that you have installed [Unity ML-Agents](#) and [NumPy](#).

```
In [2]: !pip -q install ./python
```

```
tensorflow 1.7.1 has requirement numpy>=1.13.3, but you'll have numpy 1.12.1 which is incompatible
ipython 6.5.0 has requirement prompt-toolkit<2.0.0,>=1.0.15, but you'll have prompt-toolkit 2.0.
```

```
In [3]: import numpy as np
import torch
import matplotlib.pyplot as plt
import time

from unityagents import UnityEnvironment
from collections import deque
from itertools import count
import datetime

from ddpg import DDPG, ReplayBuffer

%load_ext autoreload
%autoreload 2
%matplotlib inline
```

Next, we will start the environment! *Before running the code cell below*, change the `file_name` parameter to match the location of the Unity environment that you downloaded.

- **Mac:** "path/to/Reacher.app"

- **Windows** (x86): "path/to/Reacher\_Windows\_x86/Reacher.exe"
- **Windows** (x86\_64): "path/to/Reacher\_Windows\_x86\_64/Reacher.exe"
- **Linux** (x86): "path/to/Reacher\_Linux/Reacher.x86"
- **Linux** (x86\_64): "path/to/Reacher\_Linux/Reacher.x86\_64"
- **Linux** (x86, headless): "path/to/Reacher\_Linux\_NoVis/Reacher.x86"
- **Linux** (x86\_64, headless): "path/to/Reacher\_Linux\_NoVis/Reacher.x86\_64"

For instance, if you are using a Mac, then you downloaded Reacher.app. If this file is in the same folder as the notebook, then the line below should appear as follows:

```
env = UnityEnvironment(file_name="Reacher.app")
```

```
In [4]: #env = UnityEnvironment(file_name='envs/Reacher_Linux_NoVis_20/Reacher.x86_64') # Headless
        #env = UnityEnvironment(file_name='envs/Reacher_Linux_20/Reacher.x86_64') # Visual
        env = UnityEnvironment(file_name="/data/Reacher_Linux_NoVis/Reacher")
```

```
INFO:unityagents:
```

```
'Academy' started successfully!
```

```
Unity Academy name: Academy
```

```
    Number of Brains: 1
```

```
    Number of External Brains : 1
```

```
    Lesson number : 0
```

```
    Reset Parameters :
```

```
        goal_size -> 5.0
```

```
        goal_speed -> 1.0
```

```
Unity brain name: ReacherBrain
```

```
    Number of Visual Observations (per agent): 0
```

```
    Vector Observation space type: continuous
```

```
    Vector Observation space size (per agent): 33
```

```
    Number of stacked Vector Observation: 1
```

```
    Vector Action space type: continuous
```

```
    Vector Action space size (per agent): 4
```

```
    Vector Action descriptions: , , ,
```

Environments contain *brains* which are responsible for deciding the actions of their associated agents. Here we check for the first brain available, and set it as the default brain we will be controlling from Python.

```
In [5]: # get the default brain
        brain_name = env.brain_names[0]
        brain = env.brains[brain_name]
```

## 1.0.2 2. Examine the State and Action Spaces

In this environment, a double-jointed arm can move to target locations. A reward of +0.1 is provided for each step that the agent's hand is in the goal location. Thus, the goal of your agent is to maintain its position at the target location for as many time steps as possible.

The observation space consists of 33 variables corresponding to position, rotation, velocity, and angular velocities of the arm. Each action is a vector with four numbers, corresponding to torque applicable to two joints. Every entry in the action vector must be a number between -1 and 1.

Run the code cell below to print some information about the environment.

```
In [6]: # reset the environment
env_info = env.reset(train_mode=True)[brain_name]

# number of agents
num_agents = len(env_info.agents)
print('Number of agents:', num_agents)

# size of each action
action_size = brain.vector_action_space_size
print('Size of each action:', action_size)

# examine the state space
states = env_info.vector_observations
state_size = states.shape[1]
print('There are {} agents. Each observes a state with length: {}'.format(states.shape[0], state_size))
print('The state for the first agent looks like:', states[0])
```

Number of agents: 20

Size of each action: 4

There are 20 agents. Each observes a state with length: 33

The state for the first agent looks like: [ 0.00000000e+00 -4.00000000e+00 0.00000000e+00  
-0.00000000e+00 -0.00000000e+00 -4.37113883e-08 0.00000000e+00  
0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00  
0.00000000e+00 0.00000000e+00 -1.00000000e+01 0.00000000e+00  
1.00000000e+00 -0.00000000e+00 -0.00000000e+00 -4.37113883e-08  
0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00  
0.00000000e+00 0.00000000e+00 5.75471878e+00 -1.00000000e+00  
5.55726624e+00 0.00000000e+00 1.00000000e+00 0.00000000e+00  
-1.68164849e-01]

### 1.0.3 3. Take Random Actions in the Environment

In the next code cell, you will learn how to use the Python API to control the agent and receive feedback from the environment.

Once this cell is executed, you will watch the agent's performance, if it selects an action at random with each time step. A window should pop up that allows you to observe the agent, as it moves through the environment.

Of course, as part of the project, you'll have to change the code so that the agent is able to use its experience to gradually choose better actions when interacting with the environment!

```
In [7]: env_info = env.reset(train_mode=False)[brain_name] # reset the environment
states = env_info.vector_observations # get the current state (for each agent)
```

```

scores = np.zeros(num_agents) # initialize the score (for each
while True:
    actions = np.random.randn(num_agents, action_size) # select an action (for each agent)
    actions = np.clip(actions, -1, 1) # all actions between -1 and 1
    env_info = env.step(actions)[brain_name] # send all actions to the environment
    next_states = env_info.vector_observations # get next state (for each agent)
    rewards = env_info.rewards # get reward (for each agent)
    dones = env_info.local_done # see if episode finished
    scores += env_info.rewards # update the score (for each agent)
    states = next_states
    # roll over states to next time step
    if np.any(dones): # exit loop if episode finished
        break
    break
print('Total score (averaged over agents) this episode: {}'.format(np.mean(scores)))

```

Total score (averaged over agents) this episode: 0.0

When finished, you can close the environment.

#### 1.0.4 4. It's Your Turn!

Now it's your turn to train your own agent to solve the environment! When training the environment, set `train_mode=True`, so that the line for resetting the environment looks like the following:

```
env_info = env.reset(train_mode=True)[brain_name]
```

```
In [8]: avg_over = 100
        print_every = 10
```

```

def ddpg(agent, n_episodes=200, stopOnSolved=True):
    print('Start: ', datetime.datetime.now())
    scores_deque = deque(maxlen=avg_over)
    scores_global = []
    average_global = []
    min_global = []
    max_global = []
    best_avg = -np.inf

    tic = time.time()

    for i_episode in range(1, n_episodes+1):
        env_info = env.reset(train_mode=True)[brain_name] # reset the environment
        states = env_info.vector_observations # get the current state (f
        scores = np.zeros(num_agents) # initialize the score (f
        agent.reset()

        score_average = 0

```

```

timestep = time.time()
for t in count():
    actions = agent.act(states, add_noise=True)
    env_info = env.step(actions)[brain_name]
    next_states = env_info.vector_observations
    rewards = env_info.rewards
    dones = env_info.local_done
    agent.step(states, actions, rewards, next_states, dones)
    states = next_states
    scores += rewards
    if np.any(dones):
        break

score = np.mean(scores)
scores_deque.append(score)
score_average = np.mean(scores_deque)
scores_global.append(score)
average_global.append(score_average)
min_global.append(np.min(scores))
max_global.append(np.max(scores))
print('\r {}, {:.2f}, {:.2f}, {:.2f}, {:.2f}, {:.2f}'\
      .format(str(i_episode).zfill(3), score, score_average, np.max(scores),
              np.min(scores), time.time() - timestep), end="\n")
if i_episode % print_every == 0:
    agent.save('./')
if stopOnSolved and score_average >= 30.0:
    toc = time.time()
    print('\nSolved in {:d} episodes!\tAvg Score: {:.2f}, time: {}'.format(i_episode,
                                                                           score_average,
                                                                           toc - timestep))
    agent.save('./'+str(i_episode)+'_')
    break

print('End: ',datetime.datetime.now())
return scores_global, average_global, max_global, min_global

```

ddpg

Out[8]: <function \_\_main\_\_.ddpg(agent, n\_episodes=200, stopOnSolved=True)>

In [8]: agent = DDPG(state\_size=state\_size, action\_size=action\_size, CER=False)  
scores, averages, maxima, minima = ddpg(agent, n\_episodes=200)

```

plt.plot(np.arange(1, len(averages)+1), averages)
plt.plot(np.arange(1, len(scores)+1), scores)
plt.plot(np.arange(1, len(maxima)+1), maxima)
plt.plot(np.arange(1, len(minima)+1), minima)
plt.title('Learning without CER')
plt.ylabel('Score')
plt.xlabel('Episode #')

```

```
plt.legend(['Total Avg', 'Episode Avg', 'Max', 'Min'], loc='upper left')
plt.grid()
plt.show()
```

Start: 2019-11-19 21:38:07.228439

001,	0.48,	0.48,	2.46,	0.05,	16.86
002,	0.78,	0.63,	1.98,	0.25,	17.21
003,	0.77,	0.68,	1.62,	0.00,	17.22
004,	0.95,	0.74,	2.01,	0.51,	17.23
005,	1.13,	0.82,	2.24,	0.33,	17.33
006,	0.86,	0.83,	2.58,	0.23,	17.51
007,	1.06,	0.86,	2.67,	0.00,	17.54
008,	1.05,	0.89,	2.32,	0.12,	17.63
009,	1.23,	0.92,	2.30,	0.33,	17.77
010,	1.43,	0.97,	2.95,	0.60,	17.96
011,	1.56,	1.03,	3.41,	0.32,	18.13
012,	1.89,	1.10,	3.52,	0.48,	18.31
013,	2.03,	1.17,	3.86,	0.70,	18.68
014,	1.79,	1.22,	3.08,	0.60,	18.82
015,	2.36,	1.29,	6.42,	0.54,	19.15
016,	2.46,	1.36,	4.39,	0.92,	19.35
017,	2.01,	1.40,	4.45,	0.39,	19.72
018,	2.60,	1.47,	5.88,	1.30,	19.93
019,	3.67,	1.58,	7.28,	1.81,	20.51
020,	3.62,	1.69,	5.87,	1.92,	21.11
021,	3.60,	1.78,	7.08,	1.85,	20.85
022,	4.32,	1.89,	7.53,	2.83,	21.14
023,	4.72,	2.02,	9.00,	1.47,	21.58
024,	4.65,	2.13,	8.75,	1.94,	21.92
025,	6.19,	2.29,	13.02,	2.39,	22.28
026,	6.28,	2.44,	9.41,	2.01,	22.85
027,	6.18,	2.58,	11.42,	1.99,	23.18
028,	6.26,	2.71,	10.11,	3.60,	23.64
029,	6.95,	2.86,	9.82,	3.13,	23.89
030,	6.68,	2.99,	10.33,	3.68,	24.22
031,	7.83,	3.14,	12.92,	3.81,	24.60
032,	9.20,	3.33,	11.96,	5.27,	25.08
033,	7.42,	3.45,	10.94,	2.01,	25.74
034,	7.86,	3.58,	11.42,	3.88,	26.02
035,	8.85,	3.73,	10.74,	5.79,	26.32
036,	9.49,	3.89,	13.61,	5.48,	26.67
037,	9.86,	4.06,	18.07,	4.87,	27.01
038,	9.87,	4.21,	14.23,	4.17,	27.47
039,	10.64,	4.37,	19.29,	4.01,	27.70
040,	11.69,	4.56,	18.20,	5.40,	28.83
041,	12.41,	4.75,	15.78,	5.42,	28.59
042,	11.78,	4.92,	16.20,	9.70,	28.92
043,	11.20,	5.06,	15.65,	6.43,	29.26


044, 13.49, 5.25, 19.36, 7.14, 30.02  
045, 11.58, 5.39, 17.78, 4.85, 30.12  
046, 12.93, 5.56, 16.39, 9.15, 30.26  
047, 13.24, 5.72, 20.54, 9.84, 30.69  
048, 14.54, 5.90, 20.52, 8.78, 31.46  
049, 14.81, 6.09, 19.60, 8.28, 31.49  
050, 14.19, 6.25, 20.41, 8.36, 31.98  
051, 15.62, 6.43, 19.97, 11.01, 32.19  
052, 14.83, 6.59, 20.28, 9.72, 32.22  
053, 15.91, 6.77, 22.00, 10.61, 31.96  
054, 17.03, 6.96, 20.93, 9.36, 31.83  
055, 16.42, 7.13, 21.48, 11.24, 31.58  
056, 17.16, 7.31, 22.86, 12.99, 31.24  
057, 18.11, 7.50, 25.54, 11.36, 31.33  
058, 18.86, 7.70, 27.25, 11.33, 31.56  
059, 19.22, 7.89, 25.86, 13.86, 32.29  
060, 21.00, 8.11, 27.75, 15.41, 32.78  
061, 18.57, 8.28, 27.54, 12.15, 32.16  
062, 20.36, 8.48, 30.27, 12.30, 32.97  
063, 21.33, 8.68, 27.11, 11.87, 32.65  
064, 21.30, 8.88, 26.24, 13.86, 32.34  
065, 20.99, 9.06, 25.68, 11.34, 32.16  
066, 22.83, 9.27, 28.10, 9.52, 31.96  
067, 20.41, 9.44, 26.16, 15.13, 32.10  
068, 21.22, 9.61, 28.97, 10.12, 32.04  
069, 21.91, 9.79, 28.46, 11.80, 32.15  
070, 25.93, 10.02, 31.25, 13.45, 32.18  
071, 27.73, 10.27, 35.42, 14.51, 32.00  
072, 25.05, 10.48, 34.45, 8.54, 32.11  
073, 27.78, 10.71, 33.21, 11.48, 32.22  
074, 23.92, 10.89, 34.73, 10.61, 32.13  
075, 25.65, 11.09, 30.69, 13.46, 32.05  
076, 23.57, 11.25, 31.72, 11.13, 32.05  
077, 28.41, 11.47, 35.17, 6.56, 32.08  
078, 26.85, 11.67, 33.94, 10.23, 32.30  
079, 28.57, 11.89, 35.88, 14.93, 32.17  
080, 29.29, 12.10, 34.31, 17.38, 32.55  
081, 28.42, 12.30, 39.47, 10.88, 32.10  
082, 25.55, 12.47, 34.27, 12.29, 32.73  
083, 28.44, 12.66, 34.63, 16.47, 31.74  
084, 28.77, 12.85, 34.69, 12.65, 32.00  
085, 30.42, 13.06, 39.37, 23.16, 31.58  
086, 31.06, 13.27, 36.76, 22.66, 31.51  
087, 32.36, 13.49, 39.44, 21.62, 31.55  
088, 33.60, 13.71, 37.96, 28.97, 31.94  
089, 35.28, 13.96, 39.58, 32.71, 31.78  
090, 34.05, 14.18, 37.67, 22.87, 31.79  
091, 32.64, 14.38, 39.55, 21.06, 31.78

092, 33.78, 14.59, 39.13, 30.19, 31.81  
093, 34.62, 14.81, 39.29, 30.03, 31.90  
094, 34.17, 15.02, 38.51, 27.10, 31.72  
095, 34.07, 15.22, 37.55, 23.69, 31.98  
096, 33.59, 15.41, 39.27, 24.48, 32.02  
097, 34.81, 15.61, 38.93, 28.76, 32.08  
098, 34.79, 15.80, 39.49, 28.69, 32.13  
099, 34.07, 15.99, 39.41, 27.82, 32.17  
100, 36.19, 16.19, 39.58, 32.93, 32.75  
101, 35.89, 16.54, 39.52, 31.35, 32.26  
102, 35.36, 16.89, 39.42, 29.51, 32.41  
103, 38.01, 17.26, 39.61, 34.81, 32.35  
104, 38.23, 17.63, 39.42, 36.68, 32.08  
105, 38.39, 18.01, 39.63, 30.72, 32.24  
106, 38.27, 18.38, 39.64, 31.98, 32.14  
107, 38.26, 18.75, 39.57, 30.64, 32.43  
108, 37.93, 19.12, 39.60, 29.03, 32.25  
109, 38.42, 19.49, 39.64, 32.74, 32.20  
110, 38.81, 19.87, 39.66, 33.05, 32.14  
111, 38.67, 20.24, 39.67, 36.73, 32.25  
112, 38.86, 20.61, 39.69, 34.39, 32.43  
113, 38.17, 20.97, 39.67, 31.89, 32.25  
114, 38.98, 21.34, 39.65, 34.56, 32.21  
115, 39.25, 21.71, 39.63, 38.27, 32.17  
116, 38.23, 22.07, 39.67, 34.11, 32.06  
117, 38.06, 22.43, 39.65, 32.65, 32.33  
118, 39.34, 22.80, 39.70, 38.45, 32.30  
119, 38.89, 23.15, 39.66, 35.41, 32.10  
120, 38.86, 23.50, 39.66, 34.57, 32.87  
121, 38.72, 23.85, 39.62, 36.69, 32.20  
122, 38.69, 24.20, 39.61, 36.45, 32.57  
123, 37.94, 24.53, 39.58, 33.53, 31.67  
124, 37.66, 24.86, 39.65, 33.24, 31.69  
125, 37.22, 25.17, 39.59, 29.30, 32.19  
126, 36.98, 25.48, 39.43, 32.74, 32.35  
127, 38.26, 25.80, 39.60, 35.94, 32.15  
128, 38.43, 26.12, 39.64, 36.11, 31.36  
129, 37.97, 26.43, 39.65, 33.28, 31.48  
130, 38.37, 26.75, 39.57, 34.79, 32.12  
131, 38.99, 27.06, 39.56, 37.39, 32.37  
132, 38.28, 27.35, 39.50, 33.69, 32.37  
133, 37.85, 27.65, 39.59, 34.57, 32.34  
134, 38.46, 27.96, 39.54, 36.40, 32.28  
135, 36.51, 28.23, 39.69, 29.32, 32.24  
136, 37.63, 28.52, 39.46, 32.34, 32.20  
137, 37.43, 28.79, 39.60, 32.08, 32.12  
138, 35.49, 29.05, 39.15, 14.50, 32.23  
139, 34.04, 29.28, 39.21, 26.32, 32.24



140, 33.44, 29.50, 39.16, 26.42, 32.84  
 141, 36.43, 29.74, 39.58, 32.06, 32.54  
 142, 31.30, 29.94, 38.99, 24.32, 32.35  
 143, 35.10, 30.17, 38.88, 23.16, 32.35

Solved in 143 episodes! Avg Score: 30.17, time: 4150.681720256805  
 End: 2019-11-19 22:47:17.913165



```
In [8]: # Agent with CER enabled
agent = DDPG(state_size=state_size, action_size=action_size, CER=True)
scores, averages, maxima, minima = ddpg(agent, n_episodes=200)

plt.plot(np.arange(1, len(averages)+1), averages)
plt.plot(np.arange(1, len(scores)+1), scores)
plt.plot(np.arange(1, len(maxima)+1), maxima)
plt.plot(np.arange(1, len(minima)+1), minima)
plt.title('Learning with CER')
plt.ylabel('Score')
plt.xlabel('Episode #')
plt.legend(['Total Avg', 'Episode Avg', 'Max', 'Min'], loc='upper left')
plt.grid()
plt.show()
```

Start: 2019-11-24 18:18:30.753477

001, 0.42, 0.42, 1.23, 0.00, 20.43  
002, 0.75, 0.59, 1.43, 0.00, 20.48  
003, 0.94, 0.71, 2.52, 0.24, 20.58  
004, 0.86, 0.74, 1.62, 0.45, 20.65  
005, 1.02, 0.80, 1.97, 0.15, 20.69  
006, 1.48, 0.91, 3.02, 0.20, 20.96  
007, 1.48, 0.99, 2.37, 0.87, 21.09  
008, 1.54, 1.06, 3.00, 0.40, 21.27  
009, 1.41, 1.10, 2.98, 0.51, 21.41  
010, 1.35, 1.13, 2.86, 0.35, 21.65  
011, 1.72, 1.18, 3.38, 0.73, 21.79  
012, 2.39, 1.28, 4.07, 0.32, 21.99  
013, 2.48, 1.37, 4.21, 1.13, 22.34  
014, 3.23, 1.51, 6.44, 0.34, 22.73  
015, 3.54, 1.64, 6.32, 1.33, 22.88  
016, 3.41, 1.75, 7.19, 0.67, 23.15  
017, 4.29, 1.90, 6.57, 0.30, 23.51  
018, 4.76, 2.06, 6.51, 2.47, 23.88  
019, 5.16, 2.22, 12.07, 0.00, 24.01  
020, 5.67, 2.39, 10.73, 1.69, 24.70  
021, 6.28, 2.58, 11.05, 4.04, 24.77  
022, 7.41, 2.80, 13.95, 1.24, 25.11  
023, 10.28, 3.12, 16.78, 5.50, 25.46  
024, 12.02, 3.50, 18.92, 5.53, 25.68  
025, 14.58, 3.94, 31.15, 2.22, 26.17  
026, 21.46, 4.61, 31.63, 8.99, 26.63  
027, 27.02, 5.44, 35.90, 12.16, 27.34  
028, 23.78, 6.10, 34.21, 10.90, 27.37  
029, 26.48, 6.80, 34.00, 17.22, 27.85  
030, 25.41, 7.42, 33.92, 14.75, 28.24  
031, 31.13, 8.19, 35.64, 24.32, 28.81  
032, 32.25, 8.94, 38.52, 19.88, 29.04  
033, 33.22, 9.67, 36.32, 26.30, 29.56  
034, 36.13, 10.45, 39.48, 30.47, 29.84  
035, 35.75, 11.17, 39.45, 28.17, 30.29  
036, 38.42, 11.93, 39.58, 35.52, 31.09  
037, 36.59, 12.60, 39.68, 27.79, 31.12  
038, 36.50, 13.23, 39.20, 32.97, 31.75  
039, 37.71, 13.85, 39.67, 35.74, 31.98  
040, 37.77, 14.45, 39.61, 36.24, 32.71  
041, 37.36, 15.01, 39.65, 35.61, 32.54  
042, 37.94, 15.56, 39.64, 32.68, 32.82  
043, 37.79, 16.07, 39.68, 35.54, 33.61  
044, 38.29, 16.58, 39.22, 36.83, 34.15  
045, 37.96, 17.05, 39.61, 35.53, 34.03  
046, 39.05, 17.53, 39.54, 38.09, 34.97  
047, 38.52, 17.98, 39.66, 33.98, 35.53


048, 38.63, 18.41, 39.65, 29.94, 35.77  
049, 36.67, 18.78, 38.46, 28.65, 35.51  
050, 37.57, 19.16, 39.51, 33.42, 35.97  
051, 38.40, 19.53, 39.66, 33.28, 36.27  
052, 38.16, 19.89, 39.64, 34.51, 36.27  
053, 36.29, 20.20, 38.98, 30.51, 36.30  
054, 33.01, 20.44, 35.23, 26.88, 36.41  
055, 34.27, 20.69, 38.70, 29.47, 36.49  
056, 37.27, 20.99, 39.39, 33.18, 36.65  
057, 35.43, 21.24, 39.03, 32.46, 36.76  
058, 36.35, 21.50, 39.57, 32.50, 36.38  
059, 37.61, 21.77, 39.67, 35.49, 36.56  
060, 36.90, 22.03, 39.59, 32.53, 37.27  
061, 35.54, 22.25, 38.80, 31.98, 36.49  
062, 38.05, 22.50, 39.63, 34.77, 36.65  
063, 35.67, 22.71, 39.58, 32.27, 36.63  
064, 36.02, 22.92, 39.55, 30.39, 36.83  
065, 38.54, 23.16, 39.59, 37.48, 36.46  
066, 38.16, 23.39, 39.63, 36.42, 36.59  
067, 38.91, 23.62, 39.65, 37.97, 36.55  
068, 37.91, 23.83, 39.58, 35.62, 36.79  
069, 38.09, 24.04, 39.59, 35.45, 36.50  
070, 38.79, 24.25, 39.65, 35.15, 36.63  
071, 38.73, 24.45, 39.48, 37.43, 36.64  
072, 38.53, 24.65, 39.64, 36.21, 36.60  
073, 38.35, 24.83, 39.56, 35.70, 36.52  
074, 38.83, 25.02, 39.61, 37.09, 36.52  
075, 38.78, 25.21, 39.54, 37.02, 36.56  
076, 38.66, 25.38, 39.58, 36.95, 36.67  
077, 39.02, 25.56, 39.67, 37.10, 37.23  
078, 38.86, 25.73, 39.65, 36.32, 36.81  
079, 39.17, 25.90, 39.64, 37.53, 36.68  
080, 38.42, 26.06, 39.69, 36.35, 37.21  
081, 38.49, 26.21, 39.61, 30.36, 36.78  
082, 38.78, 26.36, 39.52, 36.88, 37.38  
083, 38.66, 26.51, 39.65, 35.37, 37.45  
084, 37.42, 26.64, 39.64, 29.81, 37.49  
085, 36.14, 26.75, 39.33, 27.88, 37.51  
086, 37.02, 26.87, 39.45, 33.27, 37.53  
087, 36.96, 26.99, 38.99, 33.80, 37.46  
088, 37.11, 27.10, 38.76, 35.11, 37.61  
089, 37.06, 27.22, 39.49, 33.46, 37.62  
090, 36.82, 27.32, 38.91, 32.59, 37.60  
091, 36.86, 27.43, 39.11, 33.80, 37.65  
092, 37.23, 27.53, 39.49, 34.30, 37.58  
093, 37.64, 27.64, 39.37, 33.32, 37.66  
094, 35.43, 27.73, 38.24, 30.13, 37.76  
095, 35.64, 27.81, 38.07, 32.89, 37.71

```

096, 32.18, 27.85, 36.93, 27.21, 37.78
097, 33.81, 27.92, 36.96, 31.19, 37.25
098, 36.56, 28.00, 39.56, 31.06, 36.88
099, 36.79, 28.09, 39.38, 27.31, 36.83
100, 37.48, 28.19, 39.20, 33.30, 37.46
101, 37.92, 28.56, 39.57, 34.51, 36.71
102, 38.82, 28.94, 39.49, 37.28, 36.75
103, 38.69, 29.32, 39.62, 35.48, 36.79
104, 36.05, 29.67, 39.67, 31.91, 36.90
105, 33.13, 29.99, 37.59, 27.24, 36.75
106, 36.73, 30.35, 39.52, 30.26, 37.05

```

Solved in 106 episodes! Avg Score: 30.35, time: 3425.3794577121735  
End: 2019-11-24 19:15:36.136079



Saves experiences for training future agents. Warning file is quite large.

```
In [24]: env.close()
```

### 1.0.5 5. See the pre-trained agent in action

```
In [10]: agent = DDPG(state_size=state_size, action_size=action_size, CER=True)
```

```
In [11]: agent.load('./96_96_actor.pth', './106_96_96_critic.pth')
```

```

In [12]: def play(agent, episodes=3):
          for i_episode in range(episodes):
              env_info = env.reset(train_mode=False)[brain_name]      # reset the environment
              states = env_info.vector_observations                    # get the current state
              scores = np.zeros(num_agents)                          # initialize the score (
              while True:
                  actions = np.random.randn(num_agents, action_size) # select an action (for
                  actions = agent.act(states, add_noise=False)        # all actions between -1
                  env_info = env.step(actions)[brain_name]           # send all actions to tn
                  next_states = env_info.vector_observations           # get next state (for ea
                  rewards = env_info.rewards                           # get reward (for each a
                  dones = env_info.local_done                          # see if episode finishe
                  scores += env_info.rewards                           # update the score (for
                  states = next_states                                 # roll over states to ne
                  if np.any(dones):                                    # exit loop if episode f
                      break
                  #break
          print('Ep No: {} Total score (averaged over agents): {}'.format(i_episode, np.m

In [13]: play(agent, 10)

```

```

Ep No: 0 Total score (averaged over agents): 34.48299922924489
Ep No: 1 Total score (averaged over agents): 31.785499289538713
Ep No: 2 Total score (averaged over agents): 32.320999277569356
Ep No: 3 Total score (averaged over agents): 33.78049924494699
Ep No: 4 Total score (averaged over agents): 31.428999297507108
Ep No: 5 Total score (averaged over agents): 32.94849926354364
Ep No: 6 Total score (averaged over agents): 31.92699928637594
Ep No: 7 Total score (averaged over agents): 32.104499282408504
Ep No: 8 Total score (averaged over agents): 33.96599924080074
Ep No: 9 Total score (averaged over agents): 32.25999927893281

```